

Elements Of Data Science - S2022

Week 3: Pandas, Data Exploration and Visualization

2/1/2022

TODOs

- **Read** Selections from PDSH Chapter 3
- **Read** Selections from PDSH Chapter 4
- (Optional) Seaborn Tutorial **<https://seaborn.pydata.org/tutorial.html>**
- Complete Week 3 Quiz
- HW1 out this week, includes questions on Hypothesis Testing

TODAY

- Pandas
- Data Exploration
- Visualization in Python

Questions?

Environment Setup

In [1]:

```
import numpy as np
```

Intro to Pandas



Pandas is an open source, BSD-licensed library providing:

- **high-performance, easy-to-use data structures** and
- **data analysis tools**

In [2]:

```
# usually imported using the alias 'pd'  
import pandas as pd
```

- Primary datastructures:
 - **Series:** 1D array with a flexible index
 - **Dataframe:** 2D matrix with flexible index and column names

Pandas Series

- 1D array of data (any numpy datatype) plus an associated **index** array

In [3]:

```
s = pd.Series(np.random.rand(4))  
s
```

Out[3]:

```
0    0.472418  
1    0.924422  
2    0.033124  
3    0.441399  
dtype: float64
```

In [4]:

```
# return the values of the series  
s.values
```

Out[4]:

```
array([0.47241843, 0.9244225 , 0.03312375, 0.44139901])
```

In [5]:

```
# return the index of the series  
s.index
```

Out[5]:

```
RangeIndex(start=0, stop=4, step=1)
```


Pandas Series Cont.

- index is flexible, can be anything hashable (integers, strings, ...)

In [6]:

```
# create Series from array and set index  
s = pd.Series([1,2,3],index=['a','b','c'],name='Example_Series')  
s
```

Out[6]:

a	1
b	2
c	3

Name: Example_Series, dtype: int64

In [7]:

```
s['a']
```

Out[7]:

1

In [8]:

```
s[['b','c']]
```

Out[8]:

b 2

c 3

Name: Example_Series, dtype: int64

Pandas Series Cont.

- accessing other Series attributes

In [9]:

```
s
```

Out[9]:

```
a    1
b    2
c    3
```

Name: Example_Series, dtype: int64

In [10]:

```
print(f'{s.index = :}')
print(f'{s.values = :}')
print(f'{s.name = :>20s}')
print(f'{s.dtype = :}')
print(f'{s.shape = :}')
'{:>20s}'.format(s.name)
```

```
s.index = Index(['a', 'b', 'c'], dtype='object')
s.values = [1 2 3]
s.name = Example_Series
```

```
s.dtype = int64  
s.shape = (3,)
```

```
Out[10]:
```

```
'      Example_Series'
```

Pandas Series Cont.

In [11]:

```
# Can create series with index from a dictionary  
s = pd.Series({'a':1, 'b':2, 'c':3, 'd':4})  
s
```

Out[11]:

```
a      1  
b      2  
c      3  
d      4  
dtype: int64
```

In [12]:

```
print(f'{s.index = :}')  
print(f'{s.values = :}')
```

```
s.index = Index(['a', 'b', 'c', 'd'], dtype='object')  
s.values = [1 2 3 4]
```


Pandas DataFrame

- tabular datastructure
- each column a single datatype
- contains both row and column indices
- single column == Series

Pandas DataFrame Cont.

In [13]:

```
df = pd.DataFrame({'Year':[2017,2018,2018,2019],  
                  'Class_Name':['A','A','B','A'],  
                  'Measure1':[2.1,3.0,2.4,1.9]  
                  })
```

In [14]:

df

Out[14]:

	Year	Class_Name	Measure1
0	2017	A	2.1
1	2018	A	3.0
2	2018	B	2.4
3	2019	A	1.9

In [15]:

```
print(df)
```

	Year	Class_Name	Measure1
0	2017	A	2.1

1	2018	A	3.0
2	2018	B	2.4
3	2019	A	1.9

In [16]:

```
display(df)
```

	Year	Class_Name	Measure1
0	2017	A	2.1
1	2018	A	3.0
2	2018	B	2.4
3	2019	A	1.9

Pandas DataFrame Cont.

In [17]:

```
data = [[2017, 'A', 2.1],  
        [2018, 'A', 3.0],  
        [2018, 'B', 2.4],  
        [2019, 'A', 1.9]]
```

In [18]:

```
df = pd.DataFrame(data,  
                  columns=['Year', 'Class_Name', 'Measure1'],  
                  index=['001', '002', '003', '004'])  
df.shape
```

Out[18]:

(4, 3)

In [19]:

```
df
```

Out[19]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0
003	2018	B	2.4

	Year	Class_Name	Measure1
004	2019	A	1.9

Pandas Attributes

- Get shape of DataFrame : `shape`

In [20]:

```
df.shape # rows, columns
```

Out[20]:

```
(4, 3)
```

- Get index values : `index`

In [21]:

```
df.index
```

Out[21]:

```
Index(['001', '002', '003', '004'], dtype='object')
```

- Get column values : `columns`

In [22]:

```
df.columns
```

Out[22]:

```
Index(['Year', 'Class_Name', 'Measure1'], dtype='object')
```

Pandas Indexing/Selection

Select by label:

- `.loc[]`

In [23]:

```
df.loc['001']
```

Out[23]:

Year	2017
Class_Name	A
Measure1	2.1

Name: 001, dtype: object

In [24]:

```
df.loc['001', 'Measure1']
```

Out[24]:

2.1

Pandas Indexing/Selection Cont.

Select by position:

- `.iloc[]`

In [25]:

```
df.iloc[0]
```

Out[25]:

Year	2017
Class_Name	A
Measure1	2.1

Name: 001, dtype: object

In [26]:

```
df.iloc[0,2]
```

Out[26]:

2.1

Pandas Indexing/Selection Cont.

Selecting multiple rows/columns: use list (fancy indexing)

In [27]:

```
df.loc[['002', '004']]
```

Out[27]:

	Year	Class_Name	Measure1
002	2018	A	3.0
004	2019	A	1.9

In [28]:

```
df.loc[['002', '004'], ['Year', 'Measure1']]
```

Out[28]:

	Year	Measure1
002	2018	3.0
004	2019	1.9

Pandas Slicing

In [29]:

```
# Get last two rows  
df.iloc[-2:]
```

Out[29]:

	Year	Class_Name	Measure1
003	2018	B	2.4
004	2019	A	1.9

In [30]:

```
# Get first two rows and first two columns  
df.iloc[:2,:2]
```

Out[30]:

	Year	Class_Name
001	2017	A
002	2018	A

NOTE: `.iloc` is **exclusive** (start:end+1)

Pandas Slicing Cont.

Can also slice using labels:

In [31]:

```
df.loc['002':'004']
```

Out[31]:

	Year	Class_Name	Measure1
002	2018	A	3.0
003	2018	B	2.4
004	2019	A	1.9

In [32]:

```
df.loc['002':'004', : 'Class_Name']
```

Out[32]:

	Year	Class_Name
002	2018	A
003	2018	B
004	2019	A

NOTE: `.loc` is **inclusive**

Pandas Slicing Cont.

How to indicate all rows or all columns? :

In [33]:

```
df.loc[:, 'Measure1']
```

Out[33]:

001 2.1

002 3.0

003 2.4

004 1.9

Name: Measure1, dtype: float64

In [34]:

```
df.iloc[2:,:] ]
```

Out[34]:

	Year	Class_Name	Measure1
003	2018	B	2.4
004	2019	A	1.9

Pandas Indexing Cont.

Shortcut for indexing:

In [35]:

```
df['Class_Name']
```

Out[35]:

001	A
-----	---

002	A
-----	---

003	B
-----	---

004	A
-----	---

Name: Class_Name, dtype: object

In [36]:

```
# can use dot notation if there is no space in label  
df.Class_Name
```

Out[36]:

001	A
-----	---

002	A
-----	---

003	B
-----	---

004 A

Name: Class_Name, dtype: object

Panda Selection Chaining

Get 'Year' and 'Measure1' for first 3 rows:

In [37]:

```
df.iloc[:3].loc[:,['Year','Measure1']]
```

Out[37]:

	Year	Measure1
001	2017	2.1
002	2018	3.0
003	2018	2.4

For records '001' and '003' get last two columns

In [38]:

```
df.loc[['001','003']].iloc[:, -2:]
```

Out[38]:

	Class_Name	Measure1
001	A	2.1
003	B	2.4

Panda Selection Chaining Cont.

For record '001' get last two columns?:

In [39]:

```
# reduce the amount of error information printed
%xmode Minimal
```

Exception reporting mode: Minimal

In [40]:

```
# Note: add 'raises-exception' tag to cell to continue running after exception
df.loc['001'].iloc[:, -2:] # row with label '001', then all rows, last two columns?
```

IndexingError: Too many indexers

In [41]:

```
df.loc['001']
```

Out[41]:

Year	2017
Class_Name	A

Measure1 2.1
Name: 001, dtype: object

In [42]:

```
df.loc['001'].iloc[-2:] # row with label '001', last two elements of Series
```

Out[42]:

Class_Name A
Measure1 2.1
Name: 001, dtype: object

Pandas `head` and `tail`

Get a quick view of the first or last rows in a DataFrame

In [43]:

```
df.head() # first 5 rows by default
```

Out[43]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0
003	2018	B	2.4
004	2019	A	1.9

In [44]:

```
df.tail(2) # only print last 2 rows
```

Out[44]:

	Year	Class_Name	Measure1
003	2018	B	2.4
004	2019	A	1.9

Pandas Boolean Mask

In [45]:

```
# Which rows have Class_Name of 'A'?  
df.loc[:, 'Class_Name'] == 'A'
```

Out[45]:

```
001      True  
002      True  
003     False  
004      True  
Name: Class_Name, dtype: bool
```

In [46]:

```
# Get all data for rows with with Class_Name 'A'  
df.loc[df.Class_Name == 'A']
```

Out[46]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0
004	2019	A	1.9

In [47]:

```
# Get Measure1 for all records for Class_Name 'A'  
df.loc[df.Class_Name == 'A', 'Measure1']
```

Out[47]:

001 2.1

002 3.0

004 1.9

Name: Measure1, dtype: float64

Pandas Boolean Mask Cont.

Get all records for class 'A' before 2019

In [48]:

```
df.loc[(df.Class_Name == 'A') & (df.Year < 2019)]
```

Out[48]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0

Get all records in a set of years:

In [49]:

```
df.loc[df.Year.isin([2017,2019])]
```

Out[49]:

	Year	Class_Name	Measure1
001	2017	A	2.1
004	2019	A	1.9

Pandas Selection Review

- `.loc[]`
- `.iloc[]`
- Fancy Indexing
- Slicing
- Chaining
- `head` and `tail`
- Boolean Mask

Pandas Sorting

In [50]:

```
df.sort_values(by=['Measure1']).head(3)
```

Out[50]:

	Year	Class_Name	Measure1
004	2019	A	1.9
001	2017	A	2.1
003	2018	B	2.4

In [51]:

```
df.sort_values(by=['Measure1'],ascending=False).head(3)
```

Out[51]:

	Year	Class_Name	Measure1
002	2018	A	3.0
003	2018	B	2.4
001	2017	A	2.1

In [52]:

```
df.sort_values(by=['Year','Measure1']).head(3)
```

Out[52]:

	Year	Class_Name	Measure1
001	2017	A	2.1
003	2018	B	2.4
002	2018	A	3.0

Questions?

Exploratory Data Analysis

For a new set of data, would like to know:

- amount of data (rows, columns)
- range (min, max)
- counts of discrete values
- central tendencies (mean, median)
- dispersion or spread (variance, IQR)
- skew
- covariance and correlation ...

Yellowcab Dataset

- Records of Yellowcab Taxi trips from January 2017
- more info: **<https://www1.nyc.gov/site/tlc/about/tlc-trip-record-data.page>**

Loading Datasets from CSV (Comma Separated Values)

- columns separated by delimiter, eg. comma, tab (\t), pipe (|)
- one row per record, observation
- often, strings quoted
- often, first row contains column headings
- often, comment rows starting with #

In [53]:

```
!head ../data/yellowcab_demo_withdaycategories.csv
```

```
# A sample of yellocab taxi trip data from Jan 2017
```

```
pickup_datetime,dropoff_datetime,trip_distance,fare_amount,tip_amount,payment_type,day_of_week,is_weekend
```

```
2017-01-05 14:49:04,2017-01-05 14:53:53,0.89,5.5,1.26,Credit card,3,True
```

2017-01-15 01:07:22,2017-01-15 01:26:47,2.7,14.0,
0.0,Cash,6,False

2017-01-29 09:55:00,2017-01-29 10:04:43,1.41,8.0,
0.0,Cash,6,False

2017-01-10 05:40:12,2017-01-10 05:42:22,0.4,4.0,0.
0,Cash,1,True

2017-01-06 17:02:48,2017-01-06 17:16:10,2.3,11.0,
0.0,Cash,4,True

2017-01-14 19:03:14,2017-01-14 19:08:41,0.8,5.5,,C
redit card,5,True

2017-01-06 18:51:52,2017-01-06 18:55:45,0.2,4.5,0.
0,Cash,4,True

2017-01-04 20:47:30,2017-01-04 21:01:24,2.68,11.

5,,Credit card,2,True

Loading Datasets with Pandas

In [54]:

```
import pandas as pd
df = pd.read_csv('../data/yellowcab_demo_withdaycategories.csv',
                  sep=',',
                  header=1,
                  parse_dates=['pickup_datetime', 'dropoff_datetime'])
```

In [55]:

```
# display first 5 rows
df.head(5)
```

Out[55]:

	pickup_datetime	dropoff_datetime	trip_distance	fare_amount	tip_amount	payment_type	day_of_week	is_weekend
0	2017-01-05 14:49:04	2017-01-05 14:53:53	0.89	5.5	1.26	Credit card	3	True
1	2017-01-15 01:07:22	2017-01-15 01:26:47	2.70	14.0	0.00	Cash	6	False
2	2017-01-29 09:55:00	2017-01-29 10:04:43	1.41	8.0	0.00	Cash	6	False
3	2017-01-10 05:40:12	2017-01-10 05:42:22	0.40	4.0	0.00	Cash	1	True

	pickup_datetime	dropoff_datetime	trip_distance	fare_amount	tip_amount	payment_type	day_of_week	is_weekend
4	2017-01-06 17:02:48	2017-01-06 17:16:10	2.30	11.0	0.00	Cash	4	True

Get Size of Dataset

In [56]:

```
df.shape
```

Out[56]:

```
(1000, 8)
```

In [57]:

```
# number of rows  
f'{df.shape[0]} rows'
```

Out[57]:

```
'1000 rows'
```

In [58]:

```
# number of columns  
f'{df.shape[1]} columns'
```

Out[58]:

```
'8 columns'
```

In [59]:

```
'number of rows: {}, number of columns: {}'.format(*df.shape)
```

Out[59]:

```
'number of rows: 1000, number of columns: 8'
```

Aside: Argument Unpacking with *

- * in when calling a function unpacks an iterable, passing each value as an argument
- want `format(2,8)` instead of the `format((2,8))`

In [60]:

```
df.shape
```

Out[60]:

```
(1000, 8)
```

In [61]:

```
# call .format( (2,8) )  
'number of rows: {}, number of columns: {}'.format(df.shape)
```

IndexError: Replacement index 1 out of range for positional args tuple

In [62]:

```
# call .format(2,8)  
'number of rows: {}, number of columns: {}'.format(*df.shape)
```

Out[62]:

'number of rows: 1000, number of columns: 8'

What are the column names?

In [63]:

```
df.columns
```

Out[63]:

```
Index(['pickup_datetime', 'dropoff_datetime', 'trip_distance', 'fare_amount',  
      'tip_amount', 'payment_type', 'day_of_week', 'is_weekend'],  
      dtype='object')
```

In [64]:

```
df.columns.values
```

Out[64]:

```
array(['pickup_datetime', 'dropoff_datetime', 'trip_distance',  
      'fare_amount', 'tip_amount', 'payment_type', 'day_of_week',  
      'is_weekend'], dtype=object)
```


In [65]:

```
df.columns.tolist()
```

Out[65]:

```
['pickup_datetime',  
 'dropoff_datetime',  
 'trip_distance',  
 'fare_amount',  
 'tip_amount',  
 'payment_type',  
 'day_of_week',  
 'is_weekend']
```

What are the column datatypes?

In [66]:

```
df.dtypes
```

Out[66]:

pickup_datetime	datetime64[ns]
dropoff_datetime	datetime64[ns]
trip_distance	float64
fare_amount	float64
tip_amount	float64
payment_type	object
day_of_week	int64
is_weekend	bool
dtype:	object

In [67]:

```
type(df.dtypes)
```

Out[67]:

```
pandas.core.series.Series
```


Get Summary Info for DataFrame

In [68]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	pickup_datetime	1000 non-null	datetime64
[ns]			
1	dropoff_datetime	1000 non-null	datetime64
[ns]			
2	trip_distance	1000 non-null	float64
3	fare_amount	1000 non-null	float64
4	tip_amount	910 non-null	float64
5	payment_type	1000 non-null	object
6	day_of_week	1000 non-null	int64

```
7    is_weekend    1000 non-null    bool  
dtypes: bool(1), datetime64[ns](2), float64(3), in  
t64(1), object(1)  
memory usage: 55.8+ KB
```

- number of rows
- number of columns
- column names, number of filled values, datatypes
- number of each datatype seen
- size of dataset in memory

Variable (Observation) Types

- **Numeric** (eg. weight, temperature)
 - usually has a zero value
 - describes magnitude
- **Categorical** (eg. class, variety)
 - usually a finite set
 - no order
- **Ordinal** (eg. Likert scale, education level, etc.)
 - usually a finite set
 - has order
 - usually missing zero
 - difference between levels may not be the same

Numeric: Data Ranges

In [69]:

```
df.trip_distance.min()
```

Out[69]:

0.0

In [70]:

```
df.trip_distance.max()
```

Out[70]:

32.77

In [71]:

```
df.min(numeric_only=True)
```

Out[71]:

trip_distance	0.0
fare_amount	2.5
tip_amount	0.0
day_of_week	0

```
is_weekend      False  
dtype: object
```

In [72]:

```
df.max(numeric_only=True)
```

Out[72]:

```
trip_distance    32.77  
fare_amount      88.0  
tip_amount       22.7  
day_of_week      6  
is_weekend       True  
dtype: object
```


Numeric: Central Tendency with Mean

- Sample Mean

$$\bar{x} = \frac{1}{n} \sum x_i$$

In [73]:

```
df.fare_amount.mean()
```

Out[73]:

12.4426

In [74]:

```
print(f'{df.fare_amount.mean() = :0.2f}')
```

df.fare_amount.mean() = 12.44

- Mean is sensitive to *outliers*
- **Outlier:** a data point that differs significantly from other observations
 - data error
 - effect of heavy tailed distribution?

Numeric: Central Tendency with Median

- Median
 - Divides sorted dataset into two equal sizes
 - 50% of the data is less than or equal to the median

In [75]:

```
df.fare_amount.median()
```

Out[75]:

9.0

- Median is *robust* to outliers
- **Robust:** Not affected by outliers

Numeric: Quantiles/Percentiles

- **Quantile::** cut point for splitting distribution
- **Percentile:** $x\%$ of data is less than or equal to the x th percentile

In [76]:

```
df.fare_amount.quantile(.95) # 95% of the data is less than or equal to x?
```

Out[76]:

33.5

In [77]:

```
df.fare_amount.quantile([.05,.95]) # 90% of the data is between 4 and 33.5
```

Out[77]:

0.05 4.0

0.95 33.5

Name: fare_amount, dtype: float64

In [78]:

```
df.fare_amount.quantile([0,.25,.5,.75,1]) # Quartiles: 25% of data is between each pair
```

Out[78]:

0.00	2.5
------	-----

0.25	6.5
------	-----

0.50	9.0
------	-----

0.75	14.0
------	------

1.00	88.0
------	------

Name: fare_amount, dtype: float64

Numeric: Spread with Variance

- Sample Variance

$$s^2 = \frac{\sum (x - \bar{x})^2}{n-1}$$

In [79]:

```
df.fare_amount.var().round(3)
```

Out[79]:

116.809

but this is in dollars²!

Numeric: Spread with Standard Deviation

- Sample Standard Deviation

$$s = \sqrt{\frac{\sum (x - \bar{x})^2}{n-1}}$$

In [80]:

```
df.fare_amount.std().round(3)
```

Out[80]:

10.808

- Back in original scale of dollars
- Sensitive to outliers

Numeric: Exploring Spread with IQR

- Quartiles
 - ~25% of data is \leq first quartile, 25th percentile
 - ~50% of data is \leq second quartile, 50th percentile (Median)
 - ~75% of data is \leq third quartile, 75th percentile
- Can find quartiles with: pandas quantile or numpy percentile
- **Interquartile Range (IQR)**
 - (third quartile - first quartile) or (75th percentile - 25th percentile)

In [81]:

```
df.fare_amount.quantile(.75) - df.fare_amount.quantile(.25)
```

Out[81]:

7.5

- IQR is robust to outliers

Numeric: Exploring Distribution with Skew

- **Skewness**

- measures asymmetry of distribution around mean
- indicates tail to left (neg) or right (pos)
- skew will lead to difference between median and mean

In [82]:

```
df.fare_amount.skew()
```

Out[82]:

2.882730031010152

Easier to understand with a plot...

Numeric Summary Stats with `.describe`

In [83]:

```
df.describe()
```

Out[83]:

	trip_distance	fare_amount	tip_amount	day_of_week
count	1000.000000	1000.000000	910.000000	1000.000000
mean	2.880010	12.442600	1.766275	2.987000
std	3.678534	10.807802	2.315507	2.043773
min	0.000000	2.500000	0.000000	0.000000
25%	0.950000	6.500000	0.000000	1.000000
50%	1.565000	9.000000	1.350000	3.000000
75%	3.100000	14.000000	2.460000	5.000000
max	32.770000	88.000000	22.700000	6.000000

Applying Functions to Groups of Data

In [84]:

```
df.groupby('payment_type')
```

Out[84]:

<pandas.core.groupby.generic.DataFrameGroupBy object at 0x7f40577e6d00>

In [85]:

```
df.groupby('payment_type').mean()
```

Out[85]:

	trip_distance	fare_amount	tip_amount	day_of_week	is_weekend
payment_type					
Cash	2.732209	11.856716	0.000000	2.898507	0.847761
Credit card	2.961870	12.761086	2.683322	3.039216	0.850679
No charge	0.500000	5.000000	0.000000	0.500000	1.000000

In [86]:

```
# applying multiple aggregation functions
df.groupby('payment_type')['trip_distance'].agg(['mean', 'median'])
```

Out[86]:

	mean	median
--	------	--------

payment_type	mean	median
Cash	2.732209	1.37
Credit card	2.961870	1.70
No charge	0.500000	0.50

In [87]:

```
df[df.payment_type.isin(['Cash', 'Credit card'])].groupby(['payment_type', 'is_weekend']).trip_distance.agg(['mean', 'median'])
```

Out[87]:

		mean	median
payment_type	is_weekend		
Cash	False	3.507059	2.10
	True	2.593063	1.28
Credit card	False	3.304646	1.74
	True	2.901702	1.70

Aside: Dealing with long chains

- long chains may not be visible in notebooks

In [88]:

```
# df[df.payment_type.isin(['Cash', 'Credit card'])].groupby(['payment_type', 'is_weekend']).trip_distance.agg(['mean', 'median'])
```

In [89]:

```
# use backslashes
df[df.payment_type.isin(['Cash', 'Credit card'])]\
    .groupby(['payment_type', 'is_weekend'])\
    .trip_distance.agg(['mean', 'median'])
```

Out[89]:

		mean	median
payment_type	is_weekend		
Cash	False	3.507059	2.10
	True	2.593063	1.28
Credit card	False	3.304646	1.74
	True	2.901702	1.70

In [90]:

```
# wrap in parentheses
(df[df.payment_type.isin(['Cash', 'Credit card'])]
 .groupby(['payment_type', 'is_weekend']))
```

```
.trip_distance.agg(['mean', 'median'])
)
```

Out[90]:

		mean	median
payment_type	is_weekend		
Cash	False	3.507059	2.10
	True	2.593063	1.28
Credit card	False	3.304646	1.74
	True	2.901702	1.70

Questions?

Visualizations in Python

- plotting with `matplotlib.pyplot`
- plotting with `pandas`
- plotting with `seaborn`
- need interactive plots? `plotly`

Matplotlib.pyplot

In [91]:


```
import matplotlib.pyplot as plt  
  
%matplotlib inline
```

In [92]:

```
plt.scatter(df.trip_distance, df.fare_amount)
```


Out[92]:

<matplotlib.collections.PathCollection at 0x7f409f77b8e0>



In [93]:


```
plt.scatter(df.trip_distance,df.fare_amount);
```



Matplotlib Axes

In [94]:


```
fig,ax = plt.subplots(1,1,figsize=(6,4))

ax.scatter(x=df.trip_distance,
           y=df.fare_amount,
           marker='x',
           color='red'
           )

ax.set_xlabel('trip_distance')
ax.set_ylabel('fare_amount')

ax.set_xlim([-10,50])
ax.set_ylim([-10,100])

ax.set_title('trip_distance vs fare_amount');
```



Matplotlib: Subplots, Figure and Axis

In [95]:

```
fig,ax = plt.subplots(1,2,figsize=(12,4))

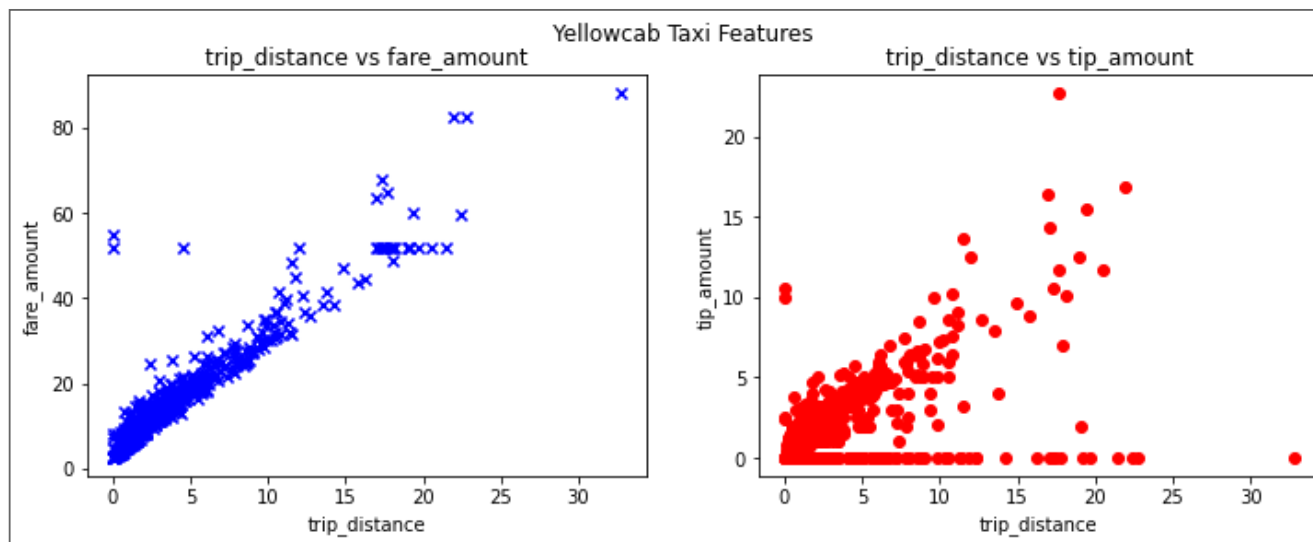
ax[0].scatter(df.trip_distance,df.fare_amount,marker='x',color='blue')
ax[1].scatter(df.trip_distance,df.tip_amount,color='red');

ax[0].set_xlabel('trip_distance')
ax[1].set_xlabel('trip_distance')

ax[0].set_ylabel('fare_amount'), ax[1].set_ylabel('tip_amount')

ax[0].set_title('trip_distance vs fare_amount')
ax[1].set_title('trip_distance vs tip_amount')


fig.suptitle('Yellowcab Taxi Features');
```



Plotting via Pandas

In [96]:


```
ax = df.plot.scatter(x='trip_distance',y='fare_amount');  
ax.set_ylabel('fare')  
ax.set_title('trip_distance vs fare_amount');
```



Univariate Distribution: Histogram


In [97]:

```
df.fare_amount.plot.hist();
```



In [98]:

```
ax = df.fare_amount.plot.hist(bins=100)  
ax.set_xlabel('fare_amount');
```



Univariate Distribution: Histogram


In [99]:

```
fig,ax = plt.subplots(1,1,figsize=(12,6));

df.fare_amount.plot.hist(bins=100, ax=ax);
ax.set_xlabel('fare_amount (dollars)');

# add a vertical line
ax.axvline(df.fare_amount.mean(),color='r');
#ax.vlines(df.fare_amount.mean(),*ax.get_ylim(),color='r');


# add some text
ax.text(df.fare_amount.mean()+1,ax.get_ylim()[1]*.75,'mean');
```



Subplots with Pandas

In [100]:

```
fig,ax = plt.subplots(1,2,figsize=(16,4))
df[df.pickup_datetime.dt.hour < 12].fare_amount.plot.hist(bins=100,ax=ax[0]);
ax[0].set_xlabel('fare_amount (dollars)');
ax[0].set_title('Trips Before Noon');
df[df.pickup_datetime.dt.hour >= 12].fare_amount.plot.hist(bins=100,ax=ax[1]);
ax[1].set_xlabel('fare_amount (seconds)');
ax[1].set_title('Trips After Noon');
```




Sharing Axes

In [101]:


```
fig, ax = plt.subplots(1, 2, figsize=(16, 4), sharey=True)

df[df.pickup_datetime.dt.hour < 12].fare_amount.plot.hist(bins=100, ax=ax[0]);
ax[0].set_xlabel('fare_amount (dollars)');
ax[0].set_title('Trips Before Noon');
df[df.pickup_datetime.dt.hour >= 12].fare_amount.plot.hist(bins=100, ax=ax[1]);
ax[1].set_xlabel('fare_amount (seconds)');
ax[1].set_title('Trips After Noon');
```



Plotting with Seaborn

- Python data visualization library
- Based on matplotlib.
- It provides a high-level interface for drawing attractive and informative statistical graphics.



In [102]:

```
import seaborn as sns
sns.__version__
```


Out[102]:

'0.11.2'

Univariate Distribution with Seaborn Histplot


In [103]:

```
fig,ax = plt.subplots(1,1,figsize=(6,3))  
sns.histplot(x='fare_amount',data=df,ax=ax);
```



In [104]:


```
fig,ax = plt.subplots(1,1,figsize=(6,3))  
sns.histplot(x=df.fare_amount,ax=ax);
```





Univariate Distribution with Seaborn Histplot Cont.

In [105]:

```
fig,ax = plt.subplots(1,1,figsize=(6,4))  
  
# many other parameters to play with  
sns.histplot(x='fare_amount',data=df,ax=ax,kde=True,stat='percent');
```



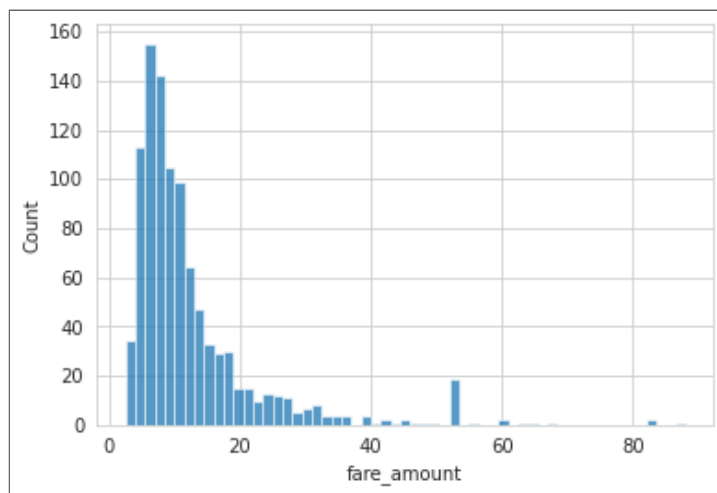
Aside: KDE



Seaborn Styles

In [106]:

```
# for a single plot using a context  
with sns.axes_style('whitegrid'):  
    sns.histplot(x=df.fare_amount);
```

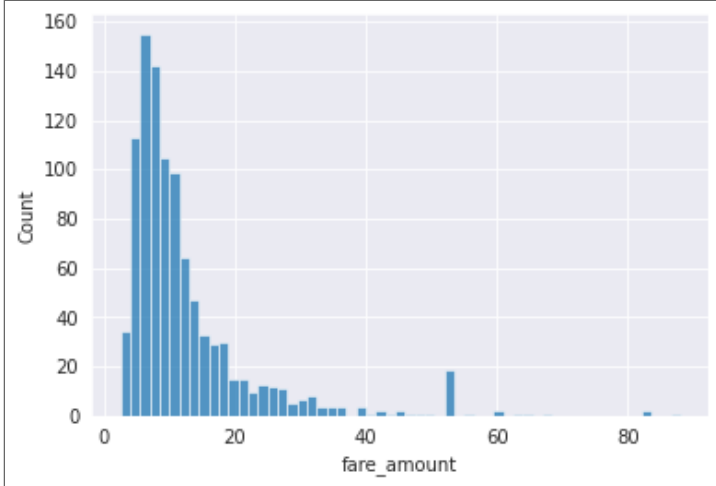


In [107]:

```
# set style globally  
sns.set_style('darkgrid')
```

In [108]:


```
sns.histplot(x=df.fare_amount);
```



Univariate Distributions: Boxplot

In [109]:


```
fig,ax = plt.subplots(1,1,figsize=(12,8))  
sns.boxplot(x=df.fare_amount,ax=ax);
```



Univariate Distributions: Boxplot

In [110]:

```
fig,ax = plt.subplots(1,1,figsize=(6,4))  
sns.boxplot(x=df.fare_amount,ax=ax);
```




- first quartile
- second quartile (Median)
- third quartile
- whiskers (usually $1.5 \times \text{IQR}$)
- outliers

Combining Plots with Subplots

In [111]:

```
fig, ax = plt.subplots(2, 1, figsize=(10, 10), sharex=True)


sns.histplot(x=df.fare_amount, ax=ax[0]);
sns.boxplot(x=df.fare_amount, ax=ax[1]);
```

Other Univariate Distribution Visualizations

In [112]:

```
fig, ax = plt.subplots(1, 3, figsize=(20, 6))  
  
sns.stripplot(x='fare_amount', data=df[:200], ax=ax[0])  
sns.violinplot(x='fare_amount', data=df, ax=ax[1])  
sns.swarmplot(x='fare_amount', data=df[:200], ax=ax[2]);
```



Bivariate: Evaluating Correlation

- **Correlation:** the degree to which two variables are linearly related
- Pearson Correlation Coefficient: $\rho_{XY} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y}$
- Sample Correlation: $r = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{(n-1)s_x s_y}$
- Takes values between:
 - -1 (highly negatively correlated)
 - 0 (not correlated)
 - 1 (highly positively correlated)

In [113]:

```
df.trip_distance.corr(df.fare_amount)
```

Out[113]:

0.948701076897808

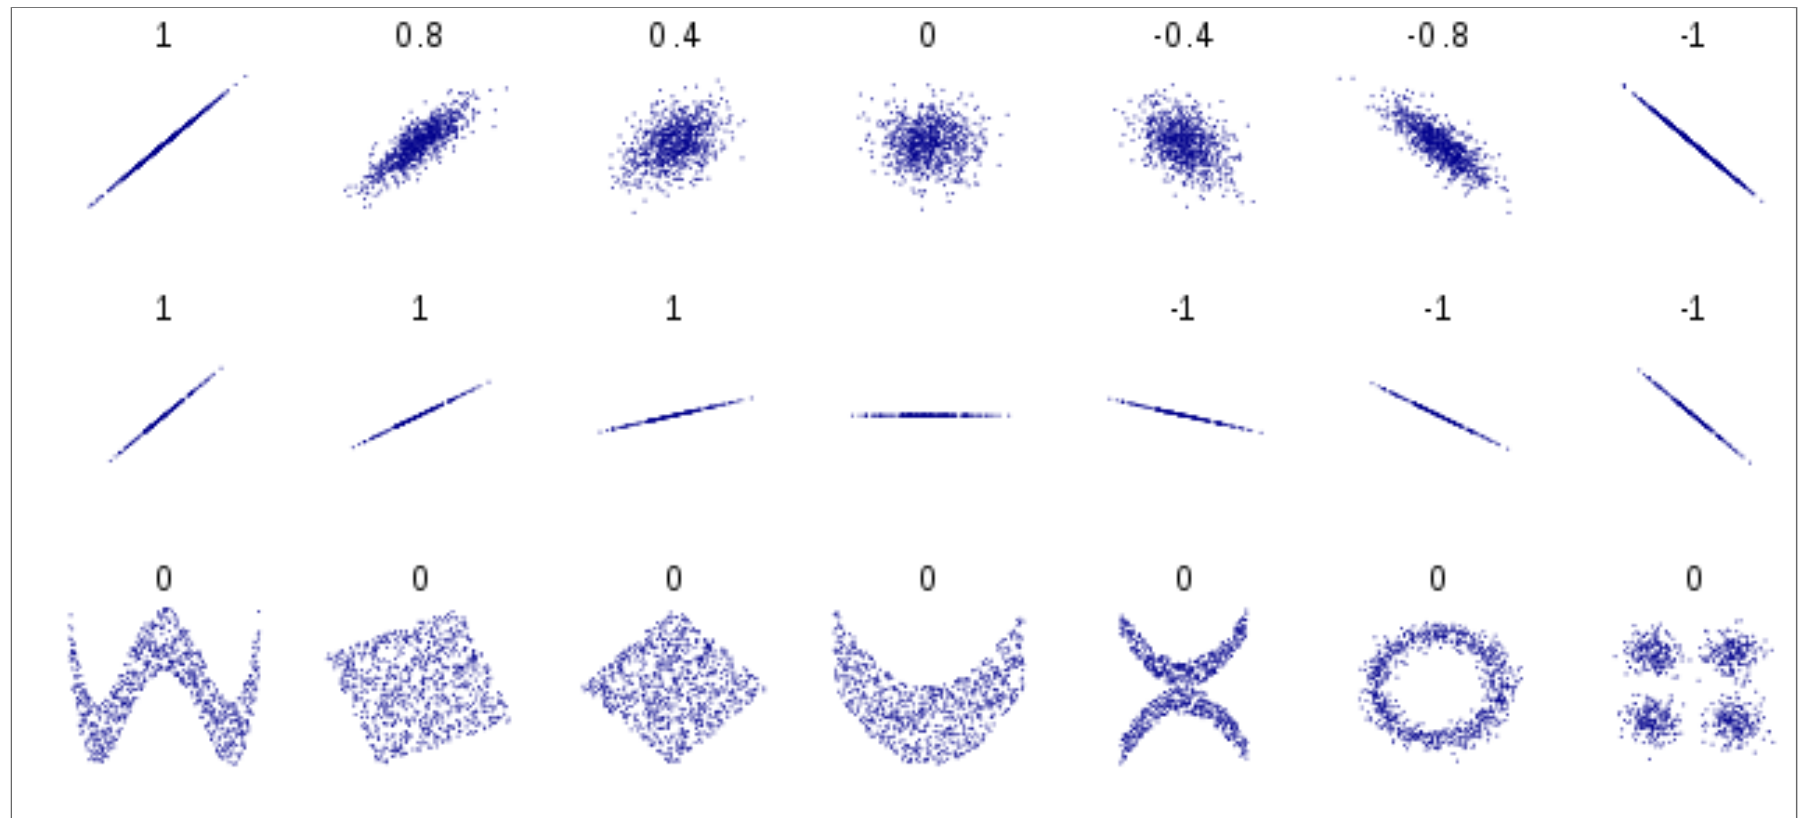
In [114]:

```
from scipy.stats import pearsonr  
r,p = pearsonr(df.trip_distance, df.fare_amount)  
r,p
```

Out[114]:

(0.9487010768978079, 0.0)

Pearson Correlation




</center>

Bivariate: Scatterplot


In [115]:

```
sns.scatterplot(x='trip_distance',y='fare_amount',data=df);
```



In [116]:


```
sns.scatterplot(x='trip_distance',y='fare_amount',data=df,alpha=0.2);
```



Bivariate: Add Regression Line

In [117]:


```
fig,ax = plt.subplots(1,1,figsize=(12,8))  
sns.regplot(x='trip_distance',y='fare_amount',data=df,ax=ax);
```



Bivariate: Joint Plot

In [118]:


```
sns.jointplot(x='trip_distance',y='fare_amount',data=df,alpha=0.5);
```



Bivariate: Joint Plot with KDE

In [119]:


```
sns.jointplot(x='trip_distance', y='fare_amount',  
              data=df,  
              kind='kde');
```



Comparing Multiple Variables with `pairplot`

In [120]:

```
sns.pairplot(df[['trip_distance', 'fare_amount', 'tip_amount']]);
```



Categorical Variables: Frequency

In [121]:

```
df.payment_type.value_counts()
```

Out[121]:

Credit card	663
Cash	335
No charge	2

Name: payment_type, dtype: int64

In [122]:

```
df.payment_type.value_counts(normalize=True)
```


Out[122]:

Credit card	0.663
Cash	0.335
No charge	0.002

Name: payment_type, dtype: float64

In [123]:


```
sns.countplot(x=df.payment_type);
```



Plotting Numeric and Categorical

In [124]:

```
fig,ax = plt.subplots(1,1,figsize=(12,8))  
  
sns.barplot(x='payment_type',y='fare_amount',data=df,estimator=np.mean,ci=95);  
ax.set_ylabel('mean fare_amount');
```




Plotting with Hue

In [125]:

```
fig,ax = plt.subplots(1,1,figsize=(12,6))


# add a second categorical variable day_of_week
sns.barplot(x='day_of_week',
            y='fare_amount',
            hue='payment_type',
            data=df);
```



Same Axis, Multiple Plots with Seaborn

In [126]:

```
fig,ax = plt.subplots(1,1,figsize=(12,6))
sns.histplot(x=df[df.pickup_datetime.dt.hour < 12].tip_amount, label='before noon',color='blue',ax=ax);
sns.histplot(x=df[df.pickup_datetime.dt.hour >= 12].tip_amount, label='after noon',color='orange',ax=ax);
plt.legend(loc='best');
```



Data Exploration and Viz Review

- central tendencies: mean, median
- spread: variance, std deviation, IQR
- correlation: pearson correlation coefficient
- plotting real valued variables: histogram, scatter, regplot
- plotting categorical variables: count, bar
- plotting interactions: jointplot, pairplot

Questions?